

Search Methods

Fundamentals of AI

Learning Outcomes

- Understand the different types of problem-solving agents
- How to formulate a problem
- Gain an in-depth understanding of basic search algorithms

A Skeletal Version of a problem-solving agent

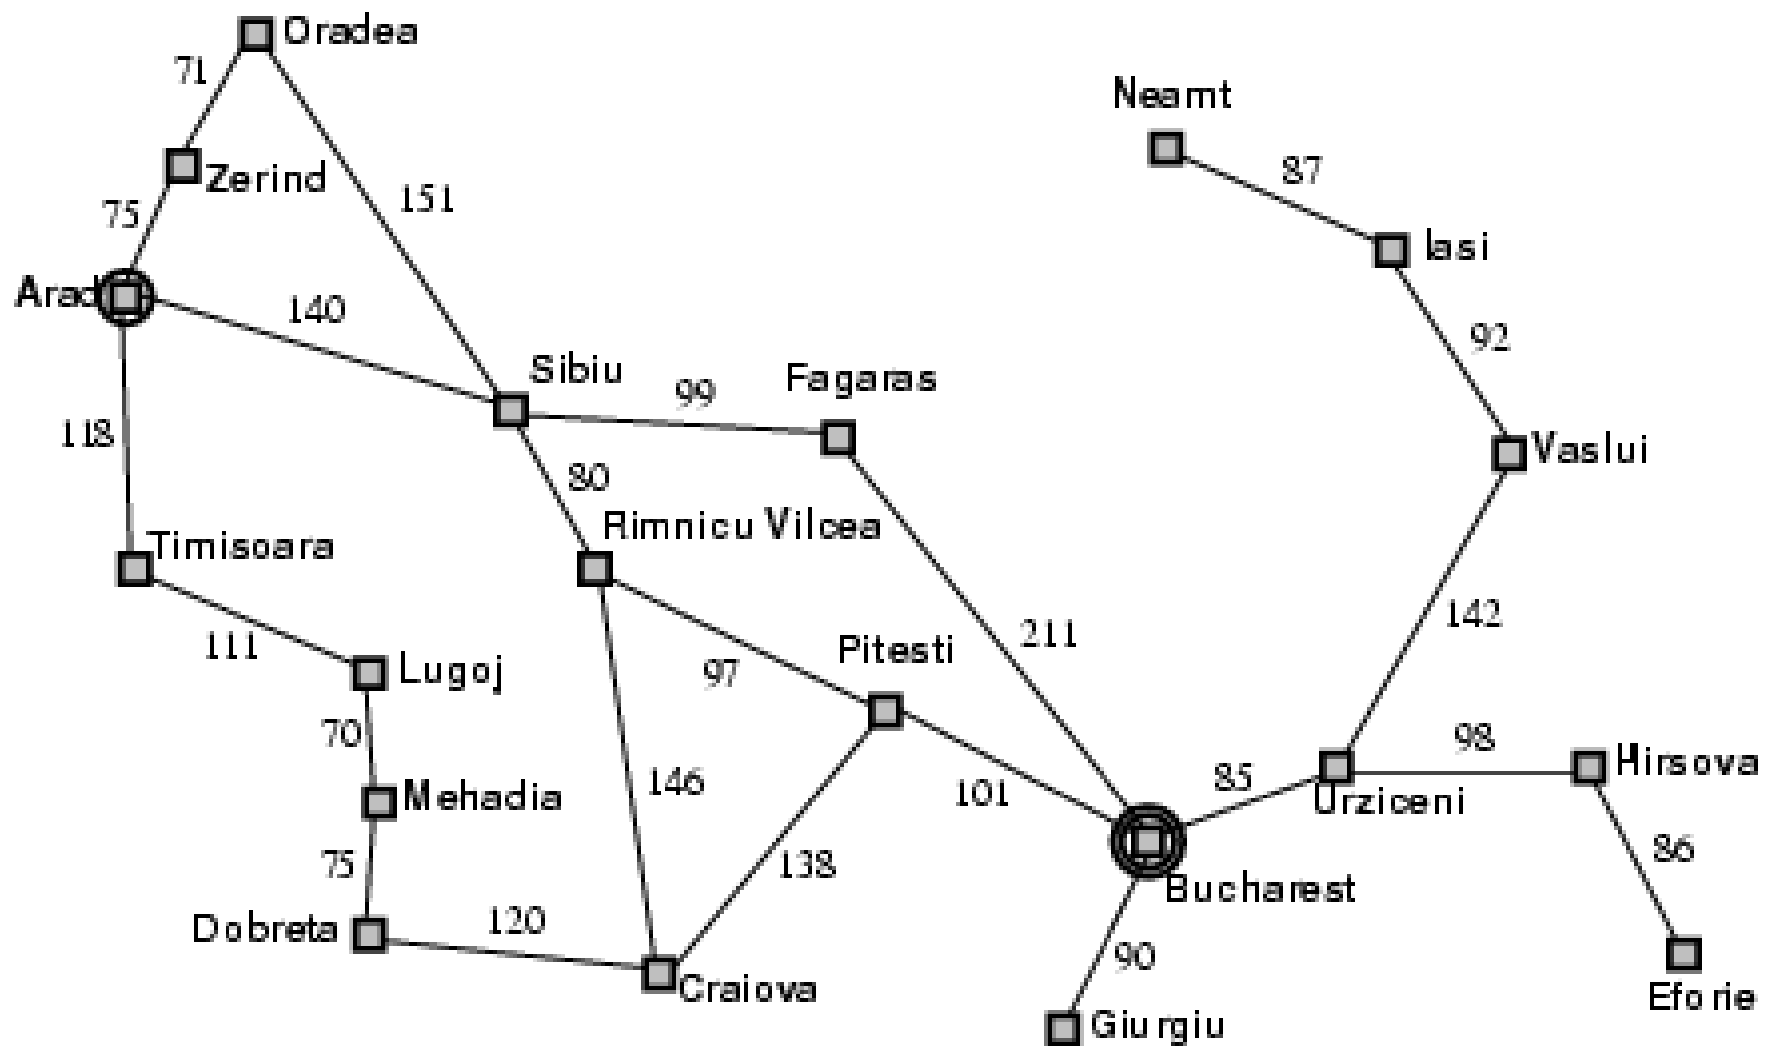
```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
           state, some description of the current world state
           goal, a goal, initially null
           problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- **Formulate goal:**
 - be in Bucharest
- **Formulate problem:**
 - **states:** various cities
 - **actions:** drive between cities
- **Find solution:**
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania




Problem types

- Deterministic, fully observable → single-state problem
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable → sensorless problem (conformant problem)
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable → contingency problem
 - percepts provide new information about current state
 - often interleave search, execution
- Unknown state space → exploration problem


Example: vacuum world

- Single-state, start in #5.
Solution?



Example: vacuum world

- **Single-state**, start in #5.
Solution? [*Right, Suck*]
- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution?



Example: vacuum world

- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}


Solution?

[Right, Suck, Left, Suck]

- **Contingency**

- Nondeterministic: *Suck* may dirty a clean carpet
- Partially observable: location, dirt at
- Percept: *[L, Clean]*, i.e., start in #5 or #7

Solution?



Example: vacuum world

- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g., *Right* goes to {2,4,6,8}


Solution?

[Right, Suck, Left, Suck]

- **Contingency**

- Nondeterministic: *Suck* may dirty a clean carpet
- Partially observable: location, dirt at cur
- Percept: *[L, Clean]*, i.e., start in #5 or #7

Solution? *[Right, **if dirt then Suck]***



Single-state problem formulation

A **problem** is defined by four items:

1. **initial state** e.g., "at Arad"
2. **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
3. **goal test**, can be
 - **explicit**, e.g., $x = \text{"at Bucharest"}$
 - **implicit**, e.g., $\text{Checkmate}(x)$
4. **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the **step cost**, assumed to be ≥ 0

A **solution** is a sequence of actions leading from the initial state to a goal state


Selecting a state space

- Real world is absurdly complex
 - state space must be **abstracted** for problem solving
 - anything irrelevant to solving the problem at hand must be left out
 - what remains is the abstract state
- Abstract state space = set of real states relevant to the problem at hand and the actions required to get to the goal state

Q) What is the abstract state space for making a black cup of coffee (without the use of a coffee machine)?


- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Vacuum world state space graph



- states?
- actions?
- goal test?
- path cost?

Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal? no dirt at all locations
- cost? 1 per action

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal?
- path?

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal? = goal state (given)
- path? 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Q) How many states exist?

Example: robotic assembly



- states?: real-valued coordinates of robot's joint angles
- actions?: continuous motions of robot joints
- goal?: complete assembly
- path?: time to complete assembly

Formulating a general solution


- Basic idea:
 - Explore the state space until the goal state is reached
 - The exploration is done according to a chosen strategy
 - The state space is structured as a tree (a natural structure)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```


Summary so far

- Problem Characteristics
 - Observable
 - Deterministic
 - Episodic
 - Static
 - Single Agent
- Problem Types
 - Single State
 - Sensorless (Conformant)
 - Contingency
 - Exploration
- Agent Types
 - Simple Reflex
 - Model-based
 - Goal-based
 - Utility-based
 - Learning
- Uninformed Search Strategies
 - Breadth-first
 - Uniform-cost
 - Depth-first
 - Depth-limited
 - Iterative Deepening
 - Bidirectional


Tree search example



Tree search example



Tree search example




Implementation: general tree search

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

```
function EXPAND(node, problem) returns a set of nodes
  successors  $\leftarrow$  the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s  $\leftarrow$  a new NODE
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
    add s to successors
  return successors
```

Implementation: states vs. nodes, what's the difference?

- A **state** is a (representation of) a physical configuration
- A **node** is part of a search tree and includes: **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `Successor Fn` of the problem to create the corresponding states.

Search strategies


- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes searched until solution is found
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

Uninformed search strategies

- **Uninformed** search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search
- Bidirectional


Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end



Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - *fringe* is a FIFO queue, i.e., new successors go at end



Properties of breadth-first search

- Complete? Yes (if b is finite)
- Time? $b + b^2 + b^3 + \dots + b^d = O(b^d)$
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal?
 - Guaranteed to find the *shallowest* goal state
 - Guaranteed to find the *least cost* goal state if path cost is a non decreasing function of depth (true when all actions have uniform cost)
- **Space** is the fundamental issue when b and d are sufficiently large

Practicalities of BFS on large scale problems

- With the processing speed taken as 1 million nodes per second, and the space required is 1kB per node (realistic assumptions) the execution times and space requirements are:

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes


Uniform-cost search

- Expands least-cost unexpanded node
- **Implementation:**
 - *fringe* = priority queue ordered by path cost
- Equivalent to BFS if step costs all equal; otherwise, is better than BFS
- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\lceil C^*/\epsilon \rceil})$ where C^* is the cost of the optimal solution
- Space? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\lceil C^*/\epsilon \rceil})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$

BFS and Uniform Cost Search: which to use?

- Uniform cost search has greater time and space complexity than that of BFS, so only use when the assumption of uniform path costs is not true
- This is generally true in path navigation problems
- Q) How about computer network routing problems?


Uniform Cost(UC) Search vs BFS: Example



- Let's say we want to get to Bucharest in Romania, starting from Sibiu
- With UC we expand the Rimnicu Vilcea (RV) node since it has the least path cost when compared to Fagaras (F), path cost so far is 80
- From RV we then expand Pitesti (P) given a path cost so far of $80+97=177$
- Now F is the node with the least path cost, expanding F gives a path cost of $99+211=310$
- We now expand P since it has a lower path cost than 310, giving the optimal path cost to Bucharest of $177+101=278$
- Q) What solution will BFS come up with?


Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO stack, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue. i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue. i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue. i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue. i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue. i.e.. put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue. i.e., put successors at front



Properties of depth-first search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if the tree is dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space!
- Optimal? No

Depth-limited search

= depth-first search with depth limit l ,
i.e., nodes at depth l have no successor

Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff  
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)  
  
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff  
  cutoff-occurred?  $\leftarrow$  false  
  if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)  
  else if DEPTH[node] = limit then return cutoff  
  else for each successor in EXPAND(node, problem) do  
    result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)  
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true  
    else if result  $\neq$  failure then return result  
  if cutoff-occurred? then return cutoff else return failure
```

Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH( problem) returns a solution, or fail-  
ure  
  inputs: problem, a problem  
  for depth  $\leftarrow$  0 to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH( problem, depth)  
    if result  $\neq$  cutoff then return result
```

Iterative deepening search / =0

Limit = 0



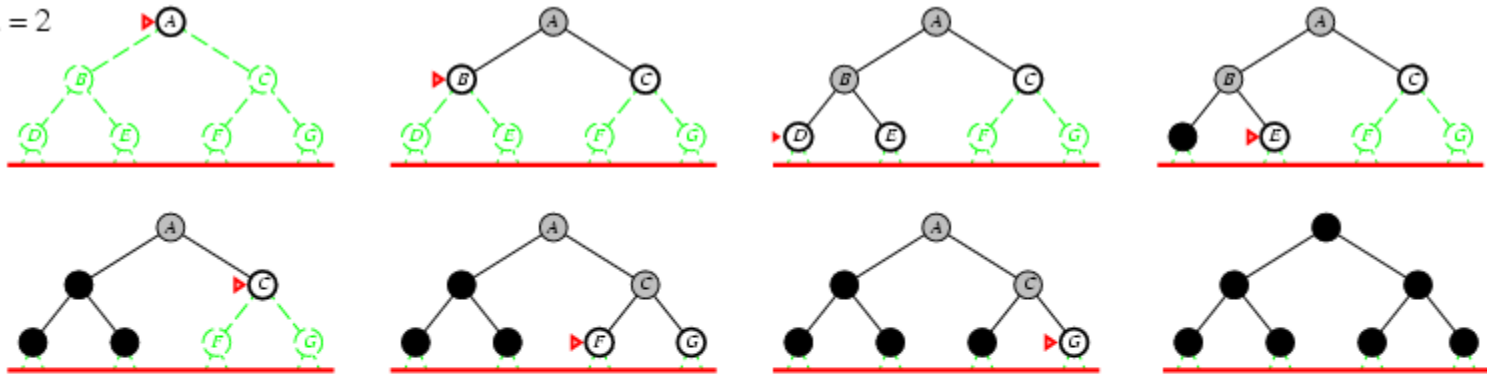
Iterative deepening search / =1

Limit = 1




Iterative deepening search $l = 2$

Limit = 2



Iterative deepening search / =3

Limit = 3



Iterative deepening search

- Number of nodes generated in a depth-limited search to depth d with branching factor b :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{l-2} + b^{l-1} + b^l$$

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For $b = 10$, $d = 5$,

- $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$

- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

- Overhead = $(123,456 - 111,111)/111,111 = 11\%$

Properties of iterative deepening search



- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1

Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Repeated states

- Failure to detect repeated states can turn a linear problem into an exponential one!



Graph search

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

Overall Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

Exam I Review (1)

- Problem Characteristics

- Observable
- Deterministic
- Episodic
- Static
- Single Agent

- Problem Types

- Single State
- Sensorless (Conformant)
- Contingency
- Exploration

- Agent Types

- Simple Reflex
- Model-based
- Goal-based
- Utility-based
- (Learning)

- Uninformed Search Strategies

- Complete?
- Optimal?
- Time
- Space

Exam I Review (2)

Uninformed Search Strategies

- Breadth-first
- Best-first
- Depth-first
- Depth-limited
- Iterative Deepening
- Bidirectional

Uniform-cost

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes